



OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLARA, CA

Fix VxLAN Issue in SFC Integration by Using Eth+NSH and VxLAN-gpe+NSH Hybrid Mode

Yi Yang, Intel (yi.y.yang@intel.com)



OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLARA, CA

Agenda

- VxLAN Issue in OVSDB+SFC
- How to Fix Current VxLAN issue by Eth+NSH
- Demo Introduction
- Demo
- Next Steps




VxLAN Issue in OVSDB + SFC + Openstack + Tacker

- The traffic to a SF(Service Function) has to be encapsulated by VxLAN-gpe with NSH header
- What will happen
 - Classifier sets tunnel ID (VNID), NSI (Network Service Index), NSP (Network Service Path) , tunnel destination IP
 - Output the packet to the VxLAN-gpe port
 - At this point, Linux kernel (vport_vxlan.ko) gets the control
 - vport_vxlan will query how to route this packet
 - If tunnel destination IP is set to 192.168.50.75, it will be routed to host2
 - If tunnel destination IP is set to 10.0.35.3, it will be dropped because the route table in host1 doesn't have an entry for that.
- The result: the packet with VxLAN-gpe + NSH header can't be routed to the SF



OPEN NET
SUMMIT 2
MARCH 14-17



```
vagrant@gbpsfc1:~$ sudo route
Kernel IP routing table
Destination      Gateway            Genmask          Flags Metric Ref  Use Iface
default        localhost          0.0.0.0          UG    0      0    0  eth0
10.0.2.0        *                255.255.255.0   U     0      0    0  eth0
172.17.0.0        *                255.255.0.0    U     0      0    0  docker0
192.168.50.0     *                255.255.255.0   U     0      0    0  eth1
vagrant@gbpsfc1:~$
```



OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLARA, CA

VxLAN Issue in OVSDB+SFC

table=0,icmp,nw_dst=10.0.36.4 actions=set_nsi:255,set_nsp:0x13,set_nshc1:0xc0a83246,set_nshc2:0x9,set_nshc3:0x12345678,
set_nshc4:0x98765432,load:0xc0a8324b->NXM_NX_TUN_IPV4_DST[],load:0x9->NXM_NX_TUN_ID[0..31],output:6

192.168.50.75

table=0,icmp,nw_dst=10.0.36.4 actions=set_nsi:255,set_nsp:0x13,set_nshc1:0xc0a83246,set_nshc2:0x9,set_nshc3:0x12345678,
set_nshc4:0x98765432,load:0x0a002303->NXM_NX_TUN_IPV4_DST[],load:0x9->NXM_NX_TUN_ID[0..31],output:6

10.0.35.3

Note: port 6 is port vxlan-gpe1




OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLARA, CA

How to Fix Current VxLAN issue by Eth + NSH

- Use Eth + NSH instead of VxLAN-gpe + NSH for communication between VMs on the same compute node/ovs bridge
- Use VxLAN-gpe + NSH only for communication between VMs on different compute nodes/ovs bridges
- Need to change ovs, openflowplugin, sfc, ovsdb, gbp to support this
 - Intel owned ovs nsh support and changed ovs to support push&pop nshð actions, but they are not committed into ovs git tree, Redhat and Intel are committed to submitting related changes into Linux kernel net subtree and ovs git tree
 - Openflowplugin changes: <https://git.opendaylight.org/gerrit/#/c/30481/>
 - Change sfcofl2 (sfc openflow renderer) and sfc classifier
 - TBD for ovsdb and gbp



OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLARA, CA





OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLAR

9.2. VXLAN-gpe + NSH

IPv4 Packet:

```
+-----+-----+
| L2 header | IP + UDP dst port=4790 |VXLAN-gpe NP=0x4(NSH)|
+-----+-----+
+-----+-----+
NSH, NP=0x1 |original packet |
+-----+-----+
```

L2 Frame:

```
+-----+-----+
| L2 header | IP + UDP dst port=4790 |VXLAN-gpe NP=0x4(NSH)|
+-----+-----+
+-----+-----+
NSH, NP=0x3 |original frame |
+-----+-----+
```

Figure 21: VXLAN-gpe + NSH

9.3. Ethernet + NSH

IPv4 Packet:

```
+-----+-----+
|Outer Ethernet, ET=0x894F | NSH, NP = 0x1 | original IP Packet |
+-----+-----+
```

L2 Frame:

```
+-----+-----+
|Outer Ethernet, ET=0x894F | NSH, NP = 0x3 | original frame |
+-----+-----+
```

Figure 22: Ethernet + NSH




OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLARA, CA


Demo Introduction

- Host1 and Host2 are vagrant boxes (VirtualBox images, Ubuntu Trusty x86_64)
- VM1, VM2, SF1 And SF2 are docker containers (Ubuntu Trusty x86_64)
- SF1 and SF2: python3 ./vxlan-tool.py --do=forward -i eth0
 - https://git.opendaylight.org/gerrit/gitweb?p=sfc.git;a=blob_plain;f=sfc-test/nsh-tools/vxlan_tool.py;h=3b05aff03aa60c556a836efc471f65d1c2595b97;hb=HEAD
- Demo two traffics
 - ICMP: ping VM2 from VM1, it will go through VM1->SF1->SF2->VM2
 - HTTP: wget <http://10.0.36.4/> in VM1, it will also go through VM1->SF1->SF2->VM2

<https://www.youtube.com/watch?v=3SIqKxZp9nY>



- 1 SFC Classifier (table 0) pushes NSH header, sfcofl2 (table 4) pushes Eth header (Destination MAC is SF1's)
- 2 sfcofl2 (table 10) pops Eth header and sets VNID and tunnel destination IP (192.168.50.75)
- 3 sfcofl2 (table 4) pushes Eth header (Destination MAC is SF2's)
- 4 sfcofl2 (table 10) pops Eth and NSH header






OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLARA, CA

Next Steps

- Commit our ovs nsh code to ovs git tree
- Push our changes in Openflowplugin (push_eth, push_nsh, pop_eth, pop_nsh) to Boron release
- Push sfc classifier and sfcofl2 changes into Boron release
- Set up the same demo in OVSDB+SFC+Openstack+Tacker integration (depends on if it is ready for such network topology)
- Push Openflowplugin Flow Programmer (A solution to openflow multi-writer issue/App coexistence issue) to Boron release
- Integrate Openflowplugin Flow Programmer to SFC, OVSDB and GBP



OPEN NETWORKING
SUMMIT 2016
MARCH 14-17, 2016 | SANTA CLARA, CA



ET: Ethernet Type
DST: Destination MAC
DTP: Destination Transmission Layer/L4 port
NP: Next Protocol
NSI: Network Service Index
 ↙ openflow action
 original packet
 NSH header
 VxLAN-gpe header
 Eth header
 IP+UDP header

Classifier flow entries:

Host1

ARP response: (ping: step 1)

table=0, priority=1000,arp,arp_tpa=10.0.35.1,arp_op=1 actions=move:NXM_OF_ETH_SRC[]->NXM_OF_ETH_DST[],
 set_field:88:f0:31:b5:12:b5->eth_src,load:0x2->NXM_OF_ARP_OP[],move:NXM_NX_ARP_SHA[]->NXM_NX_ARP_THA[],
 load:0x88f031b512b5->NXM_NX_ARP_SHA[],move:NXM_OF_ARP_SPA[]->NXM_OF_ARP_TPA[],load:0xa002301->NXM_OF_ARP_SPA[],
 IN_PORT



Routing: (ping: step 31, the second packet and later)

table=0, priority=1000,ip,dl_dst=88:f0:31:b5:12:b5,nw_dst=10.0.35.2 actions=set_field:88:f0:31:b5:12:b5->eth_src,
 set_field:00:00:00:00:35:02->eth_dst,dec_ttl,output:2

Make sure the packets from VxLAN-gpe port and the packets with Eth+NSH header can go through sfcfl2 pipeline:

table=0, priority=1001,in_port=6 actions=load:0x9->NXM_NX_REG0[],goto_table:1 (ping: step 23)
 table=0, priority=1001,encap_ether_type=20361 actions=goto_table:1 (ping: step 27)

Classifier ACL for ICMP: (ping: step 2)

table=0, priority=1000,icmp,in_port=2,nw_src=10.0.35.2,nw_dst=10.0.36.4 actions=load:0x64->NXM_NX_TUN_ID[],push_nsh,
 load:0x1->NXM_NX_NSH_MDTYPE[],load:0x3->NXM_NX_NSH_NP[],load:0x4->NXM_NX_NSP[0..23],load:0xff->NXM_NX_NS1[],
 load:0x1->NXM_NX_NSH_C1[],load:0x64->NXM_NX_NSH_C2[],load:0x3->NXM_NX_NSH_C3[],load:0x4->NXM_NX_NSH_C4[],
 goto_table:4

Classifier ACL for TCP:

table=0, priority=1000,tcp,in_port=2,nw_src=10.0.35.2,nw_dst=10.0.36.4,tp_dst=80 actions=load:0x65->NXM_NX_TUN_ID[],
 push_nsh,load:0x1->NXM_NX_NSH_MDTYPE[],load:0x3->NXM_NX_NSH_NP[],load:0x4->NXM_NX_NSP[0..23],load:0xff->NXM_NX_NS1[],
 load:0x1->NXM_NX_NSH_C1[],load:0x65->NXM_NX_NSH_C2[],load:0x3->NXM_NX_NSH_C3[],load:0x4->NXM_NX_NSH_C4[],
 goto_table:4

Classifier flow entries:

This flow entry is to make sure the packet returned back by VM2 can be delivered into VM1: (ping: step 31, the first packet)

table=0, priority=999,nsh_mdtype=0,encap_ether_type=0 actions=set_field:88:f0:31:b5:12:b5->eth_src,set_field:00:00:00:00:35:02->eth_dst, output:2

**Sfcofl2 flow entries:**

table=0, priority=5 actions=goto_table:1

Make sure the packets from VxLAN-gpe port and the packets with Eth+NSH header can go through sfcofl2 pipeline:

table=1, priority=252,reg0=0x9 actions=goto_table:4 (ping: step 24)

table=1, priority=251,encap_ether_type=20361 actions=goto_table:4 (ping: step 28)

table=1, priority=250,ip actions=goto_table:4

table=1, priority=5 actions=drop

table=2, priority=5 actions=goto_table:3

table=3, priority=5 actions=goto_table:4

push_ether for the packets from VM1, for SF1 input: (ping: step 3)

table=4, priority=550,nsi=255,nsp=4 actions=push_ether,load:0x894f->NXM_NX_ENCAP_ETHER_TYPE[],

load:0x88f031b512b5->NXM_NX_ENCAP_ETHER_SRC[],load:0x3503->NXM_NX_ENCAP_ETHER_DST[],goto_table:10

Set tunnel ID and tunnel destination IP for SF1 output:

table=4, priority=550,nsi=254,nsp=4 actions=move:NXM_NX_NSH_C2[]->NXM_NX_TUN_ID[0..31],load:0xc0a8324b->NXM_NX_TUN_IPV4_DST[],

goto_table:10

Sfc ofl2 flow entries:

push_eth for the packets from SF2, for SF1 input: (ping: step 25)

table=4, priority=550, nsi=254, nsp=8388612 actions=push_eth, load:0x894f->NXM_NX_ENCAP_ETH_TYPE[],
load:0x88f031b512b5->NXM_NX_ENCAP_ETH_SRC[], load:0x3503->NXM_NX_ENCAP_ETH_DST[], goto_table:10

table=4, priority=5 actions=goto_table:10 (ping: step 29)

pop_eth and pop_nsh and resubmit to classifier for the packets from SF1 and VM2 (Reverse RSP): (ping: step 30)

table=10, priority=660, nsi=253, nsp=8388612 actions=pop_eth, pop_nsh, resubmit(,0)

Output Eth+NSH packets from VM1 to SF1/port 3: (ping: step 4)

table=10, priority=650, nsi=255, nsp=4, encapsulate_eth_type=20361 actions=output:3

Output the Eth+NSH packets from SF2 to SF1/port 3: (ping: step 26)

table=10, priority=650, nsi=254, nsp=8388612, encapsulate_eth_type=20361 actions=output:3

pop_eth and output the packets from SF1 to VxLAN-gpe port/port 6: (ping: step 5)

table=10, priority=650, nsi=254, nsh_mdtype=1, nsh_np=3, nsp=4, encapsulate_eth_type=20361 actions=pop_eth, load:0x4->NXM_NX_TUN_GPE_NP[],
move:NXM_NX_NSH_C1[]->NXM_NX_NSH_C1[], move:NXM_NX_NSH_C2[]->NXM_NX_NSH_C2[],
move:NXM_NX_NSH_C3[]->NXM_NX_NSH_C3[], move:NXM_NX_NSH_C4[]->NXM_NX_NSH_C4[],
move:NXM_NX_TUN_IPV4_DST[]->NXM_NX_TUN_IPV4_DST[], move:NXM_NX_TUN_ID[0..31]->NXM_NX_TUN_ID[0..31], output:6

table=10, priority=5 actions=drop

Classifier flow entries:

ARP response: (ping: step 15)

```
table=0, priority=1000,arp,arp_tpa=10.0.36.1,arp_op=1 actions=move:NXM_OF_ETH_SRC[]->NXM_OF_ETH_DST[],  
set_field:88:f0:31:b5:12:b5->eth_src,load:0x2->NXM_OF_ARP_OP[],move:NXM_NX_ARP_SHA[]->NXM_NX_ARP_THA[],  
load:0x88f031b512b5->NXM_NX_ARP_SHA[],move:NXM_OF_ARP_SPA[]->NXM_OF_ARP_TPA[],load:0xa002401->NXM_OF_ARP_SPA[],IN_PORT
```

Routing: (ping: step 14, the second packet and later)

```
table=0, priority=1000,ip,dl_dst=88:f0:31:b5:12:b5,nw_dst=10.0.36.4 actions=set_field:88:f0:31:b5:12:b5->eth_src,set_field:00:00:00:00:36:04->eth_dst,dec_ttl,output:4
```

Make sure the packets from VxLAN-gpe port and the packets with Eth+NSH header can go through sfc0fl2 pipeline:

```
table=0, priority=1001,in_port=6 actions=load:0x9->NXM_NX_REG0[],goto_table:1 (ping: step 6)
```

```
table=0, priority=1001,encap_ether_type=20361 actions=goto_table:1 (ping: step 10) (ping: step 19)
```

Classifier ACL for ICMP: (ping: step 16)

```
table=0, priority=1000,icmp,in_port=4,nw_src=10.0.36.4,nw_dst=10.0.35.2 actions=load:0x66->NXM_NX_TUN_ID[],push_nsh,  
load:0x1->NXM_NX_NSH_MDTYPE[],load:0x3->NXM_NX_NSH_NP[],load:0x800004->NXM_NX_NSP[0..23],load:0xff->NXM_NX_NSI[],  
load:0x1->NXM_NX_NSH_C1[],load:0x66->NXM_NX_NSH_C2[],load:0x3->NXM_NX_NSH_C3[],load:0x4->NXM_NX_NSH_C4[],goto_table:4
```

Classifier ACL for TCP:


```
table=0, priority=1000,tcp,in_port=4,nw_src=10.0.36.4,nw_dst=10.0.35.2,tp_src=80 actions=load:0x67->NXM_NX_TUN_ID[],push_nsh,  
load:0x1->NXM_NX_NSH_MDTYPE[],load:0x3->NXM_NX_NSH_NP[],load:0x800004->NXM_NX_NSP[0..23],load:0xff->NXM_NX_NSI[],  
load:0x1->NXM_NX_NSH_C1[],load:0x67->NXM_NX_NSH_C2[],load:0x3->NXM_NX_NSH_C3[],load:0x4->NXM_NX_NSH_C4[],goto_table:4
```

This flow entry is to make sure the packets from VM1 can be delivered into VM2: (ping: step 14, the first packet)

```
table=0, priority=999,nsh_mdtype=0,encap_ether_type=0 actions=set_field:88:f0:31:b5:12:b5->eth_src,set_field:00:00:00:00:36:04->eth_dst,  
output:4
```

Sfcfl2 flow entries:

table=0, priority=5 actions=goto_table:1



OPEN NETWORKING
SOFTWARE DEVELOPMENT KIT

Make sure the packets from VxLAN-gpe port and the packets with Eth+NSH header can go through sfcfl2 pipeline:

table=1, priority=252,reg0=0x9 actions=goto_table:4 (ping: step 7)

table=1, priority=251,encap_ether_type=20361 actions=goto_table:4 (ping: step 11) (ping: step 20)

table=1, priority=250,ip actions=goto_table:4

sfcfl2 default flow entries:

table=1, priority=5 actions=drop

table=2, priority=5 actions=goto_table:3

table=3, priority=5 actions=goto_table:4

push_eth for the packets from SF1, for SF2 input: (ping: step 8)

table=4, priority=550,nsi=254,nsp=4 actions=push_eth,load:0x894f->NXM_NX_ENCAP_ETH_TYPE[],
load:0x88f031b512b5->NXM_NX_ENCAP_ETH_SRC[],load:0x3605->NXM_NX_ENCAP_ETH_DST[],goto_table:10

push_eth for the packets from VM2, for SF2 input:

table=4, priority=550,nsi=255,nsp=8388612 actions=push_eth,load:0x894f->NXM_NX_ENCAP_ETH_TYPE[],
load:0x88f031b512b5->NXM_NX_ENCAP_ETH_SRC[],load:0x3605->NXM_NX_ENCAP_ETH_DST[],goto_table:10 (ping: step 17)
table=4, priority=5 actions=goto_table:10 (ping: step 12)

Set tunnel ID and tunnel destination IP for SF2 output, for output to VxLAN-gpe port:

table=4, priority=550,nsi=254,nsp=8388612 actions=move:NXM_NX_NSH_C2[]->NXM_NX_TUN_ID[0..31],
load:0xc0a83246->NXM_NX_TUN_IPV4_DST[],goto_table:10 (ping: step 21)

pop_eth and pop_nsh and resubmit to classifier for the packets from SF1 and VM1 (Forward RSP):

table=10, priority=660,nsi=253,nsp=4 actions=pop_eth,pop_nsh,resubmit(0) (ping: step 13)

Output Eth+NSH packets from SF1/VM1 to SF2/port 5: (ping: step 9)

table=10, priority=650, nsi=254, nsp=4, encap_ether_type=20361 actions=output:5

Output Eth+NSH packets from VM2 to SF2/port 5: (ping: step 18)

table=10, nsi=255, nsp=8388612, encap_ether_type=20361 actions=output:5

pop_eth and output the packets from SF2 to VxLAN-gpe port/port 6: (ping: step 22)

table=10, priority=650, nsi=254, nsh_mdtype=1, nsh_np=3, nsp=8388612, encap_ether_type=20361 actions=pop_eth,
load:0x4->NXM_NX_TUN_GPE_NP[], move:NXM_NX_NSH_C1[]->NXM_NX_NSH_C1[],
move:NXM_NX_NSH_C2[]->NXM_NX_NSH_C2[], move:NXM_NX_NSH_C3[]->NXM_NX_NSH_C3[],
move:NXM_NX_NSH_C4[]->NXM_NX_NSH_C4[], move:NXM_NX_TUN_IPV4_DST[]->NXM_NX_TUN_IPV4_DST[],
move:NXM_NX_TUN_ID[0..31]->NXM_NX_TUN_ID[0..31], output:6

table=10, priority=5 actions=drop