

OpenVSwitch and Trace Points: Getting Started in Open Source

Mick Tarsel mjtarsel@us.ibm.com
LTC Networking Team, IBM
March 2016

Linux



Agenda


- Docker/OVS Project Setup
- Where to start
- Helpful tools & tips
- Trace events
- Tracing packets
- Closing thoughts



Linux

Project Overview

- OVS and Docker on GNU/Linux
- Sporadic latency of ARP packets in slow path



Where to Start?

- Read, a lot.
- Learn git
 - git clone <https://github.com/openvswitch/ovs>
- Don't be overwhelmed
 - Stick with small chunks
- Learn about tools

Narrowing it Down

- Knowing your environment
- cscope
 - <http://cscope.sourceforge.net/>

```
Find this C symbol: []
Find this global definition:
Find functions called by this function:
Find functions calling this function:
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files #including this file:
Find assignments to this symbol:
```

cscope screenshot



Linux

Measuring Packet Latency?

- Tcpdump?
 - I want to ID a packet
- SystemTap
 - <https://sourceware.org/systemtap/>
 - Events (upcall) and handlers (tag it)
 - Kprobes – dynamic tracing
- perf
 - Trace points – static event tracing
 - */Documentation/trace/tracepoints.txt*



Linux

Trace Points

- Record local variables & functions
- Fast tracing
- Requires some setup (not too simple)

Trace events

- <http://lwn.net/Articles/379903/>
 - “To solve this issue of automating the tracepoints, the TRACE_EVENT() macro was born.” - Rostedt
- Trace event must:
 - Create trace point
 - Create call back function
 - Record data
 - Parse data
- Output goes to /sys/kernel/debug/tracing/
- Usable from user space



Linux

My OVS Trace Event

```
TRACE_EVENT(upcall_start,  
           TP_PROTO(struct sk_buff *skb, int id),  
           TP_ARGS(skb, id),  
           TP_STRUCT__entry(  
               __field(const void *, skbaddr)  
               __field(u16, protocol)  
               __field(int, id)  
           ),  
           TP_fast_assign(  
               __entry->skbaddr = skb;  
               __entry->protocol = ntohs(skb->protocol);  
               __entry->id = (const int) id;  
           ),  
           TP_printk("skbaddr=%p proto=0x%x id=%d\n",  
                    __entry->skbaddr, __entry->protocol, __entry->id)  
);
```



Linux

Calling the Trace Event

```
int ovs_dp_upcall(struct datapath *dp, struct sk_buff *skb,
                   const struct sw_flow_key *key,
                   const struct dp_upcall_info *upcall_info)
{
    struct dp_stats_percpu *stats;
    int err;

    if (upcall_info->portid == 0) {
        err = -ENOTCONN;
        goto err;
    }

    if (trace_upcall_start_enabled())
        trace_upcall_start(skb, tag_packet(skb));


    if (!skb_is_gso(skb))
        err = queue_userspace_packet(dp, skb, key, upcall_info);
    else
        err = queue_gso_packets(dp, skb, key, upcall_info);
    if (err)
        goto err;

    return 0;
}
```



Linux

Following Packets



Output

```
# perf record -e openvswitch:upcall_start -e openvswitch:upcall_end -a
```

==== ARP Packets ===

	id	Time in Nano Seconds
Max:	5	4721362.0
Min:	6	64742.0
Average:	883686.666667	
Total ARP Packets:	6	

==== IP Packets ===

	id	Time in Nano Seconds
Max:	3	95160.0
Min:	8	44850.0

Average: 69619.0

Total IP Packets: 4

==== All Captured Packets ===

	id	Time in Nano Seconds	Protocol
Max:	5	4721362.0	0x806
Min:	8	44850.0	0x800

Average: 558059.6


Total Packets captured: 10



Results - ARP

Time:	COUNT	MIN	MAX	AVERAGE
RTT	1914	0	2001	1058
K-2-OVS	1914	0	1995	1013
OVS-2-K	1914	0	914	45

base-2 logarithmic histogram of round trip times



Count = number of packets

Value = time in Jiffies




Results - IPv4

Time:	COUNT	MIN	MAX	AVERAGE
RTT	230	0	53	5
K-2-OVS	230	0	44	2
OVS-2-K	230	0	50	2

base-2 logarithmic histogram of round trip times

value |----- count



Count = number of packets

Value = time in Jiffies



Linux

Future Plans

- Extend packet tracing tools to other projects
- ovs-benchmark tool

Acknowledgments

- IBM LTC
- LTC Networking Team
 - Pradeep Satyanarayana
 - Dave Wilder



Linux

Closing thoughts

- These are things that helped me....
 - Stay organized with git
 - READ
 - Don't be intimidated
 - Subscribe to mailing lists and answer questions
 - Use tools

Questions?


Thank you!



Linux

Socket Buffer

- `sk_buff`, `sk_buffer`, `skb`



Source image: http://vger.kernel.org/~davem/skb_data.html



Linux

tag_packet()

```
static inline int tag_packet(struct sk_buff *skb)
{
    struct packet_tag pkt_id;
    static unsigned int id;
    unsigned char *temp_pkt;

    pkt_id.eyecatch = TAGGED;
    pkt_id.reqid = ++id;

    if (sizeof(struct packet_tag) <= skb_tailroom(skb)) {
        temp_pkt = skb_put(skb, sizeof(pkt_id));
        memcpy(temp_pkt, &pkt_id, sizeof(pkt_id));
        pr_debug("%s:skb=%p reqid=%lu\n", __func__, skb, pkt_id.reqid);
        return id;
    }

    if (net_ratelimit())
        pr_debug("%s:Insufficient room to tag skb=%p\n", __func__, skb);

    return 0;
}
```



report_packet()

```
static inline int report_packet(struct sk_buff *skb)
{
    struct packet_tag *pkt_id;
    unsigned char *temp_pkt;

    temp_pkt = (skb->data)+(skb->len) - sizeof(struct packet_tag);
    pkt_id = (struct packet_tag *)temp_pkt;

    if (pkt_id->eyeCatch == TAGGED) {
        pr_debug("%s:skb=%p reqid=%lu\n", __func__, skb, pkt_id->reqid);
        return pkt_id->reqid;
    }

    return 0;
}
```



upcall_end trace point

```
/* define from upcall_start trace event a new event with the same proto,
   and args instead named upcall_end */
DEFINE_EVENT(upcall_start, upcall_end,
            TP_PROTO(struct sk_buff *skb, int id),
            TP_ARGS(skb, id)
);
```



Linux